



# Develop Generative AI Solutions with Azure OpenAI Service [AI-050]

© Copyright Microsoft Corporation. All rights reserved.

Date

1

## About this course

### What we'll cover

In this course, you will learn how to implement generative AI solutions by using Microsoft Azure OpenAI Service. Specifically, you'll learn how to:

- Provision Azure OpenAI resources and deploy OpenAI models for generative AI tasks
- Use APIs and SDKs to consume models from client applications
- Apply prompt engineering to generate relevant results, generate code and images, and use your own data

### Intended audience

The primary audience for this course is application developers seeking to include Azure OpenAI models in their applications.

Although Azure OpenAI models operate on natural language prompts, to use the APIs and SDKs effectively, a basic knowledge of Microsoft C# or Python is recommended.

© Copyright Microsoft Corporation. All rights reserved.

2

## Course schedule

### Trainers:

Find the proposed schedule in the Trainer Prep Guide (PDF) and customized if needed. Delete this note before publishing.

| Day 1           |  |
|-----------------|--|
| Early morning   | Introduction<br>Get started with Azure OpenAI Service        |
| Break           |  |
| Late morning    | Build natural language solutions<br>Apply prompt engineering |
| Lunch           |  |
| Early afternoon | Generate code<br>Generate images                             |
| Break           |  |
| Late afternoon  | Use your own data<br>Q & A                                   |

© Copyright Microsoft Corporation. All rights reserved.

3

## Get the most out of your Microsoft Learn profile

Verify, track, and share your training and certification progress and accomplishments—*all on one platform*

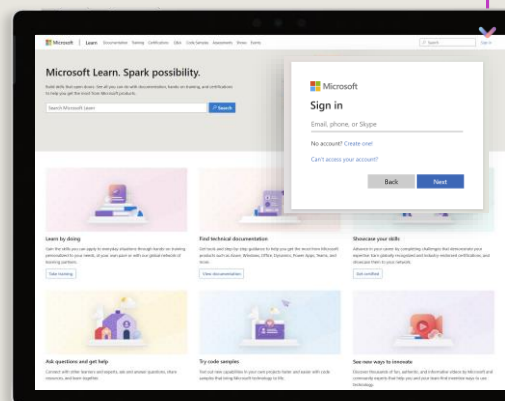
- Claim your achievement code for this course and share you have completed it.
- Access your course material and track progress on your learning activities.
- Share and verify your Microsoft Certifications via email, on social networking platforms, and on your résumé.
- Download and print transcripts and certificates.
- Manage your upcoming activities and certification exam appointments.

[www.aka.ms/MyMicrosoftLearnProfile](https://www.aka.ms/MyMicrosoftLearnProfile)

© Copyright Microsoft Corporation. All rights reserved.

### Create your Microsoft Learn profile at [learn.microsoft.com](https://learn.microsoft.com)

- Select **Sign in** at the top, right corner of any Microsoft Learn page.
- Follow the Microsoft account authentication process.
- If the account that you have chosen to sign-in with doesn't already have a Microsoft Learn profile, you'll be guided to create one.



4

## Access your course material

All course content is available on Microsoft Learn

[learn.microsoft.com/training/courses/browse](https://learn.microsoft.com/training/courses/browse)

- We'll go through this content together and as the course progresses, I will advise you on which modules to review.
- You can provide feedback for modules on Microsoft Learn. Find how at the bottom of each page.

Need help? See our [troubleshooting guide](#) or provide specific feedback by [reporting an issue](#).

© Copyright Microsoft Corporation. All rights reserved.

This course includes labs:

- Detailed lab instructions are included in your lab environment.
- Each exercise is standalone and requires:
  - A Microsoft Azure subscription
  - Approved access to the Azure OpenAI service

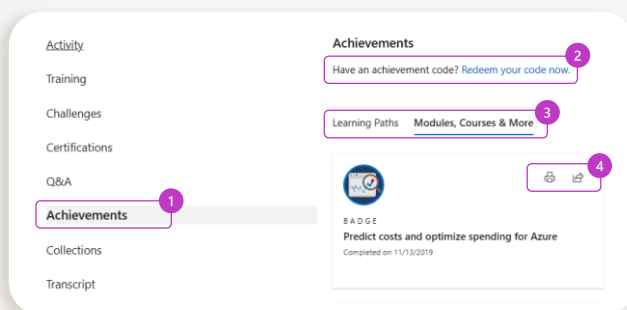
Access to the Azure OpenAI service can be requested through the form at [aka.ms/oaiapply](https://aka.ms/oaiapply)

5

## Celebrate your accomplishments and feel empowered

Get your achievement recognized and earn your badge for completing this course

[aka.ms/MyMicrosoftLearnProfile](https://aka.ms/MyMicrosoftLearnProfile)



© Copyright Microsoft Corporation. All rights reserved.

1. Go to **Achievements** on your Microsoft Learn profile
2. Redeem the code provided by your trainer
3. Find your new badge **on Modules, Courses & More**
4. Share your new achievement with your professional network. You can also download and print your certificate.

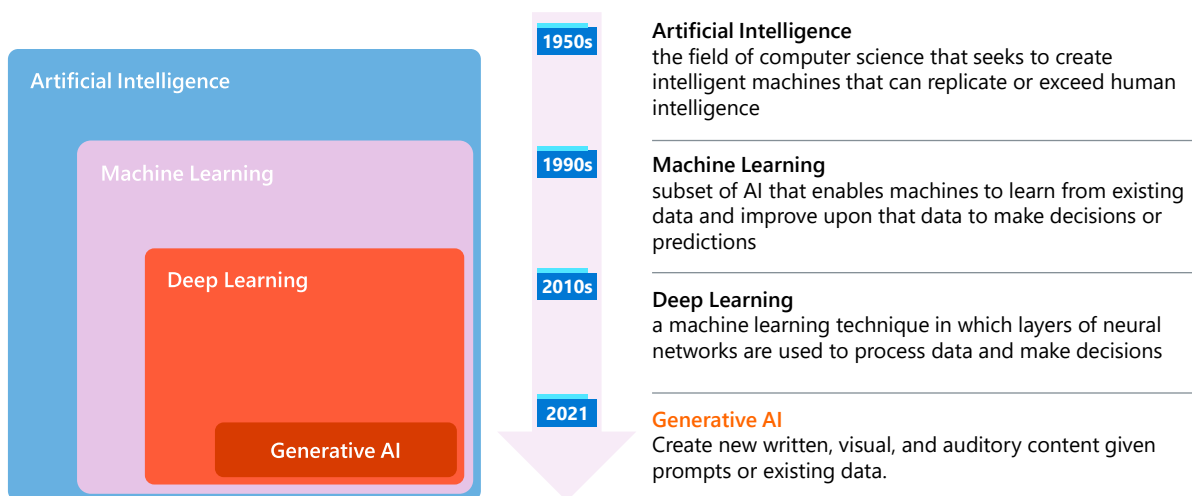
6

# Get started with Azure OpenAI Service

© Copyright Microsoft Corporation. All rights reserved.

7

## What is generative AI?



© Copyright Microsoft Corporation. All rights reserved.

8

## Provision an Azure OpenAI resource in Azure

### Deploy a model in Azure OpenAI Studio to use it

1. Apply for access to the Azure OpenAI service:  
<https://aka.ms/oaiapply>
2. Create an **Azure OpenAI** resource in the Azure portal

### Alternatively, use the Azure CLI

```
az cognitiveservices account create \
-n MyOpenAIResource \
-g MyResourceGroup \
-l eastus \
--kind OpenAI \
--sku s0 \
--subscription subscriptionID
```

Home > Azure AI services | Azure OpenAI >

### Create Azure OpenAI

1 Basics 2 Network 3 Tags 4 Review + submit

Enable new business solutions with OpenAI's language generation capabilities powered by GPT-3 models. These models have been pretrained with trillions of words and can easily adapt to your scenario with a few short examples provided at inference. Apply them to numerous scenarios, from summarization to content and code generation.

[Learn more](#)

#### Project Details

Subscription \*

Resource group \*

[Create new](#)

#### Instance Details

Region

Name \*

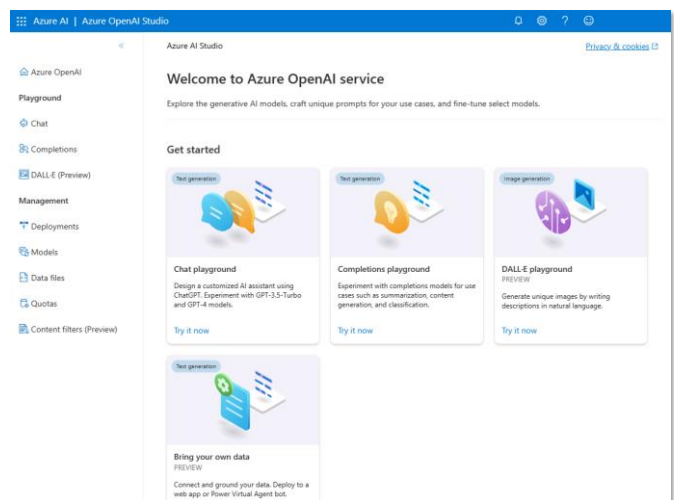
Pricing tier \*

© Copyright Microsoft Corporation. All rights reserved.

9

## Azure OpenAI Studio

- Web portal for working with Azure OpenAI models:  
<https://oai.azure.com/>
- View and deploy base models
- Connect your own data source
- Manage fine tuning and data files for custom models
- Test models in visual playgrounds:
  - **Completions** (GPT-3 and earlier models)
  - **Chat** (GPT-3.5-Turbo and later models)

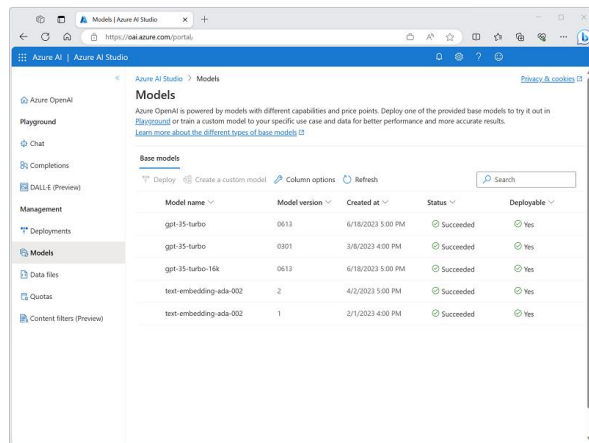


© Copyright Microsoft Corporation. All rights reserved.

10

## Types of generative AI model

| Model Family | Description   |
|--------------|---|
| GPT-4        | Newest, most capable chat-based models for language and code generation ( <i>restricted</i> ) |
| GPT-3.5      | Natural language and code-generation models   |
| Embeddings   | Models that use embeddings for specific tasks (similarity, text search, and code search)      |
| DALL-E       | Image-generation model ( <i>restricted preview</i> )  |



© Copyright Microsoft Corporation. All rights reserved.

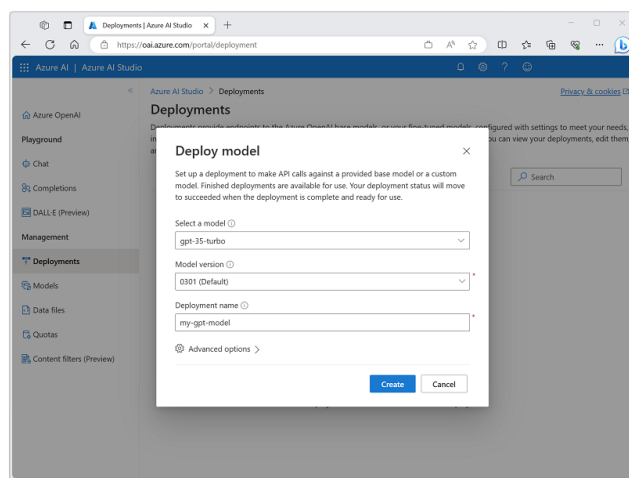
11

## Deploying generative AI models

### Deploy a model in Azure OpenAI Studio to use it

- You can deploy one or more instances of each available model
- The number of deployments depends on your quota, which you can see in the portal
- Alternatively, use the Azure CLI

```
az cognitiveservices account deployment create \
-g myResourceGroupName \
-n MyOpenAIResource \
--deployment-name my-gpt-model \
--model-name gpt-35-turbo \
--model-version "0301" \
--model-format OpenAI \
--scale-settings-scale-type "Standard"
```



© Copyright Microsoft Corporation. All rights reserved.

12

Using prompts to get completions from models

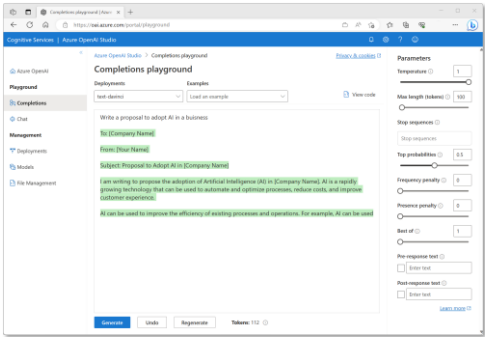
| Task                       | Prompt   | Completion   |
|----------------------------|--|--|
| Classifying content        | Tweet: I enjoyed the training course.<br>Sentiment:                              | Positive   |
| Generating new content     | Write a poem about databases   | Databases, oh databases,<br>You keep our information safe,<br>From the small to the large,<br>You store our data in a place. |
| Transformation/Translation | English: Hello<br>French:  | Bonjour  |
| Summarization              | Scotland is [long description of Scotland...]<br><br>Summarize the previous text | Scotland is [summarized description...]  |
| Continuation               | One way to grow tomatoes is to   | start with seeds...  |
| Question answering         | How many moons does Earth have?  | Earth has one moon.  |
| Chat                       | Setup, followed by messages...   | A sequence of relevant responses   |

© Copyright Microsoft Corporation. All rights reserved.

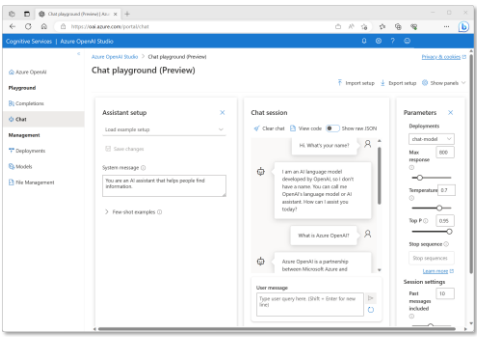
13

Testing models in Azure OpenAI Studio playgrounds

Completions playground



Chat playground



- Azure OpenAI Studio includes multiple playgrounds for specific types of model. For natural language generation models, the available playgrounds are:
- The Completions playground – used for content generation tasks with GTP-3 family models
  - The Chat playground – used for chat interactions with GPT-35-Turbo and later models.

© Copyright Microsoft Corporation. All rights reserved.

14

## Exercise: Get started with Azure OpenAI Service



Use the hosted lab environment if provided, or view the lab instructions at the link below:

<https://aka.ms/mslearn-get-started-azure-openai>

© Copyright Microsoft Corporation. All rights reserved.

15

## Resources

### Get started with Azure OpenAI Service

<https://aka.ms/mslearn-start-azure-openai>

© Copyright Microsoft Corporation. All rights reserved.

16



# Build natural language solutions with Azure OpenAI Service



17

## Integrating Azure OpenAI into your app

Applications submit prompts to models. Responses are completions.

Three REST API endpoints:

- **Completion** - model takes an input prompt, and generates one or more predicted completions
- **Embeddings** - model takes input and returns a vector representation of that input
- **ChatCompletion** - model takes input in the form of a chat conversation (where roles are specified with the message they send), and the next chat completion is generated

**ChatCompletion** will be the endpoint we focus on for this course

Use **Completion** and **Embeddings** with GPT-3 based models  
Use **ChatCompletion** with GPT-3.5-Turbo and later models

© Copyright Microsoft Corporation. All rights reserved.

18

## Using the Azure OpenAI REST API

### Completion Endpoint

`https://endpoint.openai.azure.com/openai/deployments/deployment/completions`

### Embedding Endpoint

`https://endpoint.openai.azure.com/openai/deployments/deployment/embeddings`

### ChatCompletion Endpoint

`https://endpoint.openai.azure.com/openai/deployments/deployment/chat/completions`

© Copyright Microsoft Corporation. All rights reserved.

19

## Using Azure OpenAI SDKs – Microsoft C#

```
using Azure.AI.OpenAI;

OpenAIClient client = new OpenAIClient(new Uri(yourEndpoint), new AzureKeyCredential(yourKey));

// Build completion options object
ChatCompletionsOptions chatCompletionsOptions = new ChatCompletionsOptions()
{
    Messages =
    {
        new ChatMessage(ChatRole.System, "You are a helpful assistant."),
        new ChatMessage(ChatRole.User, "What is Azure OpenAI?"),
    },
    MaxTokens = 120,
    Temperature = 0.7f,
    DeploymentName = oaiModelName
};
// Send request to Azure OpenAI model
ChatCompletions response = client.GetChatCompletions(chatCompletionsOptions);
string completion = response.Choices[0].Message.Content;
Console.WriteLine($"Answer: {completion}");
```

© Copyright Microsoft Corporation. All rights reserved.

20

## Using Azure OpenAI SDKs – Python

```
from openai import AzureOpenAI

# Initialize the Azure OpenAI client
client = AzureOpenAI(
    azure_endpoint = your_endpoint,
    api_key=your_key,
    api_version="2023-05-15")

# Send request to Azure OpenAI model
response = client.chat.completions.create(
    model=azure_oai_model,
    temperature=0.7,
    max_tokens=120,
    messages=[
        {"role": "system", "content": "You are a helpful assistant."},
        {"role": "user", "content": "Summarize the following text in 20 words or less:\n" + text}
    ]
)

print("Summary: " + response.choices[0].message.content + "\n")
```

© Copyright Microsoft Corporation. All rights reserved.

21

### Exercise: Integrate Azure OpenAI into your app



Use the hosted lab environment if provided, or view the lab instructions at the link below:

<https://aka.ms/mslearn-build-language-azure-openai>

© Copyright Microsoft Corporation. All rights reserved.

22


## Resources

Build natural language solution with Azure OpenAI Service

<https://aka.ms/mslearn-build-openai>

© Copyright Microsoft Corporation. All rights reserved.

23

A decorative graphic in the bottom right corner of the slide. It features several 3D-rendered keys in blue, purple, and red, connected by a white, winding line that suggests a path or a sequence of steps.

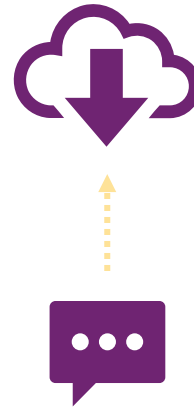
Apply prompt engineering  
with Azure OpenAI Service

24

## What is Prompt Engineering?

Constructing prompts to:

- Maximize relevancy and accuracy of completions
- Specify formatting and style of completions
- Provide conversational context
- Mitigate bias and improve fairness



© Copyright Microsoft Corporation. All rights reserved.

25

## Considerations for API endpoints

Completions API (*GPT-35-Turbo and earlier*)

Include context in the prompt (as a *meta prompt*)

```
{
  prompt: "You are a professional events planner. Write an invitation for a party to celebrate the launch of a new product."
}
```

ChatCompletions API (*GPT-35-Turbo and later*)

Set context in a system message

```
{
  "messages": [
    { "role": "system", "content": "You are a professional events planner." },
    { "role": "user", "content": "Write an invitation for a party to celebrate the launch of a new product." }
  ]
}
```

© Copyright Microsoft Corporation. All rights reserved.

26

## Providing clear instructions

Write a product description for a new water bottle



Introducing the latest addition to our product line - the innovative and eco-friendly water bottle. Made from high-quality, BPA-free materials, this bottle is designed to keep your beverages at the perfect temperature for hours. Its sleek and stylish design makes it the perfect accessory for any outdoor activity or workout session. With its convenient carrying handle and leak-proof lid, you can take it with you wherever you go, without worrying about spills or leaks. Plus, its easy-to-clean design ensures that you can use it again and again, making it the perfect choice for anyone looking for a reliable and sustainable water bottle.

Write a product description for a new water bottle that is 100% recycled. Be sure to include that it comes in natural colors with no dyes, and each purchase removes 10 pounds of plastic from our oceans



Introducing our newest water bottle, the eco-friendly choice for the conscious consumer. Made from 100% recycled materials and available in natural colors with no dyes, this bottle is not only stylish but also sustainable. Each purchase of this bottle removes 10 pounds of plastic from our oceans, making it a responsible choice for those who care about the environment. Its durable and leak-proof design makes it perfect for outdoor activities, while its easy-to-carry handle makes it convenient to take on the go. Join us in our mission to reduce plastic waste and choose this eco-friendly water bottle for your daily hydration needs.

© Copyright Microsoft Corporation. All rights reserved.

27

## Primary, supporting, and grounding content

Section markers  
(--- or ###)

Primary content to be  
summarized, translated, etc.

Supporting content to provide  
clarity / specificity

Grounding content to define  
scope for questions

---  
Thanks for your interest in our summer school program.  
We're looking for your ideas for webinars on topics such as AI, web development, ... <more...>  
---  
Topics I'm very interested in: AI, webinar dates, submission deadlines  
Extract the key points from the above email, and put them in a bulleted list  
---  
Artificial Intelligence (AI) has evolved over many years ... <more...>  
---  
Where and when did the field of AI start?

© Copyright Microsoft Corporation. All rights reserved.

28

## Cues

- Cues provide a starting point on which the completion builds
- Use cues to set the scope and format for the output
- Cues are particularly useful when generating code. For example:
  - Provide the first keyword in a SQL SELECT statement
  - Define a name and parameters for a Python function

```

---
Review: The writing was good but the movie was too long.
Review: The special effects were terrible.
<more>...
---

Summarize the reviews above.
Most common complaints:
-

```



- The movie was too long.
- The special effects were terrible.

© Copyright Microsoft Corporation. All rights reserved.

29

## Requesting output composition

Write a table in markdown with 6 animals in it, with their genus and species



| Animal           | Genus     | Species        |
|------------------|-----------|----------------|
| Red Panda        | Ailurus   | fulgens        |
| African Elephant | Loxodonta | africana       |
| Snow Leopard     | Panthera  | uncia          |
| Arctic Fox       | Vulpes    | lagopus        |
| Chimpanzee       | Pan       | troglydytes    |
| Siberian Tiger   | Panthera  | tigris altaica |

© Copyright Microsoft Corporation. All rights reserved.

30

## Using a system message

When using the ChatCompletions API, a system message sets context that can affect completions

```
{
  "role": "system",
  "content": "You are a casual, helpful assistant. You will talk like an American old western film character."
}, {
  "role": "user",
  "content": "Can you direct me to the library?"
}
```



```
{
  "role": "assistant",
  "content": "Well howdy there, stranger! The library, huh? Y'all just head down the main road till you hit the town square. Once you're there, take a left and follow the street for a spell. You'll see the library on your right, can't miss it. Happy trails!"
}
```

© Copyright Microsoft Corporation. All rights reserved.

31

## Conversation history and *few shot* learning

- Including previous messages helps retain the context and style of a conversation
- Starting a conversation with user-defined messages and responses (*few shot* learning) sets examples for the expected responses

```
{
  "role": "system",
  "content": "You are an assistant that evaluates the sentiment of customer feedback."
}, {
  "role": "user",
  "content": "That was an awesome experience"
}, {
  "role": "assistant",
  "content": "positive"
}, {
  "role": "user",
  "content": "I won't do that again"
}, {
  "role": "assistant",
  "content": "negative"
}, {
  "role": "user",
  "content": "That was not worth my time"
}, {
  "role": "assistant",
  "content": "negative"
}, {
  "role": "user",
  "content": "You can't miss this"
}
```



```
{
  "role": "assistant",
  "content": "positive"
},
```

© Copyright Microsoft Corporation. All rights reserved.

32



## Chain of thought

Ask a model to break down its response and explain its reasoning

**What sport is easiest to learn but hardest to master? Give a step by step approach of your thoughts, ending in your answer**

Step 1: Identify the criteria for "easy to learn" and "hard to master" sports.

For a sport to be considered easy to learn, it should have simple rules and require minimal equipment. For a sport to be considered hard to master, it should require years of practice to perfect and have a large variety of techniques and strategies...

Step 2: Consider different sports that fit these criteria.

Some sports that could be considered easy to learn but hard to master include:

- Tennis: Simple rules, easy to pick up, but requires years of practice to master all the different shots and strategies.
- Golf: Basic swing mechanics are easy to learn but perfecting them takes a lifetime of practice.
- ...

Step 3: Evaluate each sport based on the criteria.

Tennis and golf both require expensive equipment, which could make them less accessible to beginners. In terms of skill ceiling and time required to master, tennis and golf both have a high skill ceiling and require years of practice to perfect....

Step 4: Make a decision based on the evaluation.

Based on the above criteria and evaluation, I would say that tennis is the sport that is easiest to learn but hardest to master. While it may require expensive equipment, it is still accessible to many people. The basic rules and techniques are easy to learn, but mastering all the different shots and strategies takes years of practice...

© Copyright Microsoft Corporation. All rights reserved.

33

## Exercise: Utilize prompt engineering in your app



Use the hosted lab environment if provided, or view the lab instructions at the link below:

<https://aka.ms/openai-prompt-engineering-lab>

© Copyright Microsoft Corporation. All rights reserved.

34


## Resources

Apply prompt engineering with Azure OpenAI Service

<https://aka.ms/mslearn-prompt-engineering-openai>

© Copyright Microsoft Corporation. All rights reserved.

35

A decorative graphic in the bottom right corner of the slide, showing several 3D-rendered keyboard keys in blue, purple, and red, connected by a white line.

# Generate code with Azure OpenAI Service

36

## Use natural language prompts to generate code

- Use prompt engineering to describe what you want clearly
- Break down complex tasks into smaller pieces

### Prompt:

write a function for binary search in python

```
def binary_search(array, target):
    low = 0
    high = len(array) - 1

    while low <= high:
        mid = (low + high) // 2
        if array[mid] == target:
            return mid
        elif array[mid] < target:
            low = mid + 1
        else:
            high = mid - 1

    return -1
```

© Copyright Microsoft Corporation. All rights reserved.

37

## Change programming language and understand unknown code

### Change programming language

Convert the following code to C#:

```
def print_squares(n):
    for i in range(1, n+1):
        print(i**2)
```



```
public void PrintSquares(int n)
{
    for (int i = 1; i <= n; i++)
    {
        Console.WriteLine(i * i);
    }
}
```

© Copyright Microsoft Corporation. All rights reserved.

### Understand unknown code

Could you explain what this code is doing?:

```
fn qzplv(n: i32, m: i32) -> i32 {
    if n == 0 {
        return m + 1;
    } else if m == 0 {
        return qzplv(n - 1, 1);
    } else {
        return qzplv(n - 1, qzplv(n, m - 1));
    }
}
```



Sure, I'd be happy to explain what this code is doing!

This code is implementing the Ackermann function,...

38

## Complete partial code

- Writing your prompt with the most amount of information available will be the most accurate
- If enough code is given to determine the programming language, the model will use that language

### Prompt:

complete the following function  
"""

```
# calculate the average of the numbers
# in an array, but only if they're even
def
```

```
# calculate the average of the even
# numbers in an array
def calculate_even_average(numbers):
    total = 0
    count = 0
    for number in numbers:
        if number % 2 == 0:
            total += number
            count += 1
    if count == 0:
        return 0 # avoid division by zero
    else:
        return total / count
```

© Copyright Microsoft Corporation. All rights reserved.

39

## Write unit tests

### Prompt:

write three unit tests for this function  
"""

```
def binary_search(arr, target):
    left, right = 0, len(arr) - 1
    while left <= right:
        mid = (left + right) // 2
        if arr[mid] == target:
            return mid
        elif arr[mid] < target:
            left = mid + 1
        else:
            right = mid - 1
    return -1
```

```
def test_binary_search():
    # Test case 1: Target element is in the middle of the
    # array
    arr1 = [1, 2, 3, 4, 5, 6]
    target1 = 4
    assert binary_search(arr1, target1) == 3

    # Test case 2: Target element is at the beginning of
    # the array
    arr2 = [1, 2, 3, 4, 5, 6]
    target2 = 1
    assert binary_search(arr2, target2) == 0

    # Test case 3: Target element is not in the array
    arr3 = [1, 2, 3, 4, 5, 6]
    target3 = 7
    assert binary_search(arr3, target3) == -1
```

© Copyright Microsoft Corporation. All rights reserved.

40

## Fix bugs in your code

- Help identify and fix bugs quickly
- Provide an explanation of the issues

### Prompt:

This function isn't working for me. Fix the bugs in this function

```
"""
```

```
def calculate_average(numbers):
    total = 0
    for i in range(len(numbers)):
        number = numbers[i]
        total += number
    average = total
    rerun averave
```

```
def calculate_average(numbers):
    total = 0
    for i in range(len(numbers)):
        number = numbers[i]
        total += number
    average = total / len(numbers)
    return average
```

There were a few issues with the original code:

1. The variable "average" was assigned the value of "total" instead of the actual average.
2. There was a typo in the last line - "rerun" should be "return".
3. The average was not actually calculated - the corrected code divides the total by the length of the "numbers" list to get the average.

© Copyright Microsoft Corporation. All rights reserved.

41

## Improve performance

- Identify where improvements can be made
- Models able to optimize and refactor

### Prompt:

Can this function be more efficient?

```
"""
```

```
def sum_of_n(n):
    result = 0
    for i in range(1, n+1):
        result += i
    return result
```

We can improve this function it by using a mathematical formula instead.

Here's the more efficient version:

```
def sum_of_n(n):
    return n * (n + 1) // 2
```

© Copyright Microsoft Corporation. All rights reserved.

42

## Exercise: Generate and improve code with Azure OpenAI Service



Use the hosted lab environment if provided, or view the lab instructions at the link below:

<https://aka.ms/mslearn-openai-code-gen>

© Copyright Microsoft Corporation. All rights reserved.

43

## Resources

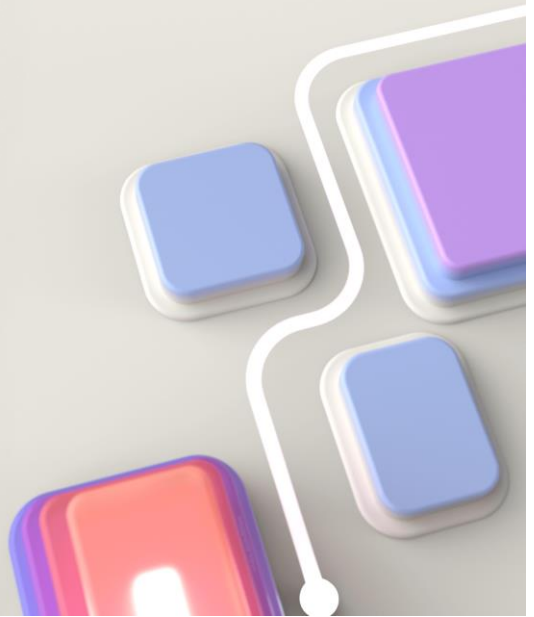
### Generate code with Azure OpenAI Service

<https://aka.ms/mslearn-code-openai>

© Copyright Microsoft Corporation. All rights reserved.

44

# Generate images with Azure OpenAI Service



45

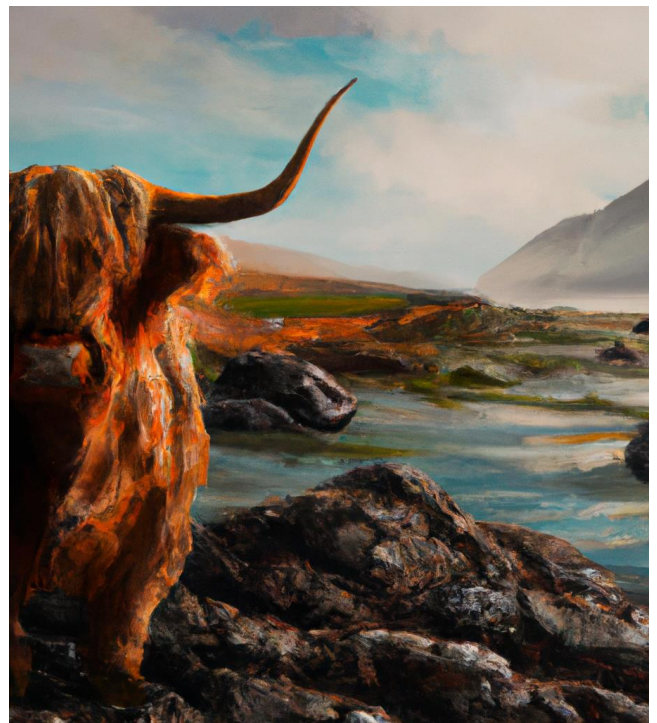
## What is DALL-E?

### Generate images with a description

- Neural network based model for generating images
- Use natural language to describe what the image should be
- Specify content and style

#### **Prompt:**

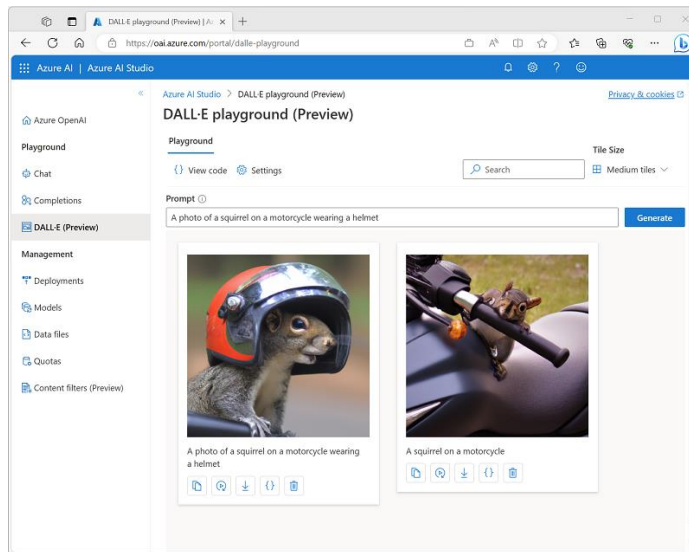
a highland cow in a field on the coast of Scotland, digital art style



© Copyright Microsoft Corporation. All rights reserved.

46

## Using DALL-E in Azure OpenAI Studio



© Copyright Microsoft Corporation. All rights reserved.

47

## Using the Azure OpenAI REST API DALL-E Endpoint

`https://endpoint.openai.azure.com/openai/images/generations:submit?api-version=api_version`

© Copyright Microsoft Corporation. All rights reserved.

48



## Using Azure OpenAI SDKs for DALL-E – Microsoft C#

```
using Azure.AI.OpenAI;

OpenAIClient client = new OpenAIClient(new Uri(yourEndpoint), new AzureKeyCredential(yourKey));

ImageGenerations response = client.GetImageGenerations(
    new ImageGenerationOptions()
    {
        Prompt = "A badger wearing a tuxedo",
        Size = ImageSize.Size1024x1024,
    });

Uri imageUri = response.Data[0].Url;
Console.WriteLine("Image URL: " + imageUri + "\n");
```

© Copyright Microsoft Corporation. All rights reserved.

49

## Using Azure OpenAI SDKs – Python

```
import openai

openai.api_type = "azure"
openai.api_base = your_endpoint
openai.api_version = "2023-06-01-preview" # be sure to use a version that supports dall-e
openai.api_key = your_key

response = openai.Image.create(
    prompt: 'A badger wearing a tuxedo',
    size: '1024x1024',
    n: 1
)

print("...Image generation complete...\n")
image_url = response["data"][0]["url"]
print("Image URL: " + image_url + "\n")
```

© Copyright Microsoft Corporation. All rights reserved.

50

## Exercise - Generate images with a DALL-E model



Use the hosted lab environment if provided, or view the lab instructions at the link below:

<https://aka.ms/mslearn-openai-dall-e>

© Copyright Microsoft Corporation. All rights reserved.

51

## Resources

Build natural language solution with Azure OpenAI Service

<https://aka.ms/mslearn-image-gen-module>

© Copyright Microsoft Corporation. All rights reserved.

52

# Use your own data with Azure OpenAI Service



53

## How Azure OpenAI can use your data



### Set up your data source

- Use an existing data source, such as an Azure search resource
- Use the Azure OpenAI studio to create that data source, if you don't already have one
- When creating the data source, you can use data already in your account such as blob storage



### Configure the studio or your app to connect to that data source

- In the studio, set up the connection by pointing it to the data source
- In your app, specify the data source in the prompt parameters



### Use the Azure OpenAI model, which now uses your data for grounding

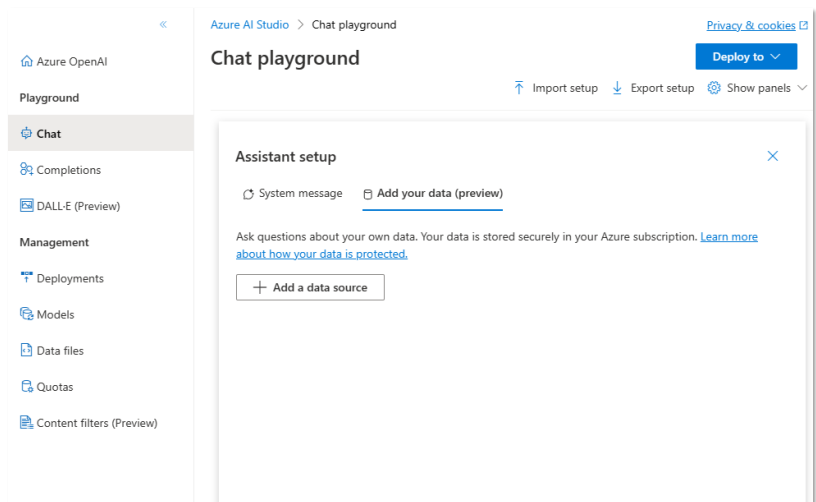
- Chat with the AI models like normal
- If the data source has relevant information about the prompt, it will use that data
- You can specify if the AI model is limited to just your data source

© Copyright Microsoft Corporation. All rights reserved.

54

## Connect to your data source

- Add your data source in the Chat playground, under Assistant setup
- Use an existing data source, or use that wizard to create a new one
- Once connected, a new chat session will start. Chat like normal, and see how the AI model references that data



© Copyright Microsoft Corporation. All rights reserved.

55

## Using the Azure OpenAI REST API

### Using your own data

<https://endpoint.openai.azure.com/openai/deployments/deployment/extensions/chat/completions?api-version=version>

- With each call, you need to specify the data source values, along with the messages array and any other parameters
- The REST endpoint will be different than normal chats

```
{
  "dataSources": [
    {
      "type": "AzureCognitiveSearch",
      "parameters": {
        "endpoint": "<your_search_endpoint>",
        "key": "<your_search_endpoint>",
        "indexName": "<your_search_index>"
      }
    }
  ],
  "messages": [
    ...
  ]
}
```

© Copyright Microsoft Corporation. All rights reserved.

56

## Using Azure OpenAI SDKs – Microsoft C#

```
AzureCognitiveSearchChatExtensionConfiguration ownDataConfig = new()
{
    SearchEndpoint = new Uri(searchEndpoint),
    IndexName = searchIndex
};
ownDataConfig.SetSearchKey(searchKey);

ChatCompletionsOptions chatCompletionsOptions = new ChatCompletionsOptions()
{
    Messages =
    {
        new ChatMessage(ChatRole.User, "I want to go to New York. Where should I stay?")
    },
    MaxTokens = 600,
    Temperature = 0.9f,
    DeploymentName = modelName,
    AzureExtensionsOptions = new AzureChatExtensionsOptions()
    {
        Extensions = {ownDataConfig}
    }
};
```

© Copyright Microsoft Corporation. All rights reserved.

57

## Using Azure OpenAI SDKs – Python

```
# Create extension config for own data
extension_config = dict(dataSources = [
    {
        "type": "AzureCognitiveSearch",
        "parameters": {
            "endpoint": search_endpoint,
            "key": search_key,
            "indexName": search_index,
        }
    }
])
response = client.chat.completions.create(
    model = azure_oai_model,
    temperature = 0.5,
    max_tokens = 1000,
    messages = [
        {"role": "user", "content": "I want to go to New York. Where should I stay?"}
    ],
    extra_body = extension_config
)
```

© Copyright Microsoft Corporation. All rights reserved.

58

## Exercise: Use your own data with Azure OpenAI Service



Use the hosted lab environment if provided, or view the lab instructions at the link below:

<https://aka.ms/mslearn-openai-own-data>

© Copyright Microsoft Corporation. All rights reserved.

59

## Resources

### Use your own data with Azure OpenAI Service

<https://aka.ms/mslearn-openai-own-data-module>

© Copyright Microsoft Corporation. All rights reserved.

60



# Thank You

© Copyright Microsoft Corporation. All rights reserved.